

Recall, Precision, and F1 Score

▼ Task

For this task, we will use the assignment from week 3 i.e. MNIST (without the addition task). Also, to make the interpretation of the scores easier, we will turn it into a binary classification problem: is the output class/digit even (0,2,4,6,8) or odd (1,3,5,7,9)

```
!pip install tableprint
```

```
Requirement already satisfied: tableprint in /usr/local/lib/python3.7/dist-packages
Requirement already satisfied: future in /usr/local/lib/python3.7/dist-packages
Requirement already satisfied: wcwidth in /usr/local/lib/python3.7/dist-packages
```

```
def score(targets, predictions):
    tp, tn, fp, fn = 0, 0, 0, 0
    for target, prediction in zip(targets, predictions):
        if target == 1:
            if prediction == 1:
                tp += 1
            else:
                fn += 1
        else:
            if prediction == 1:
                fp += 1
            else:
                tn += 1
    return tp, tn, fp, fn
```

▼ Sanity Check

After tabulating the result, we see that the calculating accuracy using our score function matches the results from our previous work (also tabulated!)

▼ Imports

```
import torch
import torch.nn as nn
import torchvision
import torchvision.transforms as transforms
import numpy as np
```

```
import torch.nn.functional as F
from torchsummary import summary
```

```
import tableprint as tp
```

▼ Loading the MNIST Dataset

```
train_dataset = torchvision.datasets.MNIST(
    root='.',
    train=True,
    transform=transforms.ToTensor(),
    download=True)
```

```
test_dataset = torchvision.datasets.MNIST(
    root='.',
    train=False,
    transform=transforms.ToTensor(),
    download=True)
```

```
/usr/local/lib/python3.7/dist-packages/torchvision/datasets/mnist.py:498: Use
return torch.from_numpy(parsed.astype(m[2], copy=False)).view(*s)
```

▼ Turn targets into binary class

```
def ev(x):
    if x%2 == 0:
        return 1
    else:
        return 0
```

```
train_dataset.targets.apply_(ev)
test_dataset.targets.apply_(ev)
```

```
tensor([0, 1, 0, ..., 1, 0, 1])
```

```
print(train_dataset.targets.max(), train_dataset.targets.min())
print(test_dataset.targets.max(), test_dataset.targets.min())
```

```
tensor(1) tensor(0)
tensor(1) tensor(0)
```

▼ Re-creating Dataset

```
# Flattens the MNIST images
train_x = train_dataset.data.reshape(60000, 784).float()
test_x = test_dataset.data.reshape(10000, 784).float()
```

```
# Creates the Dataset with the modified targets
train_ds = torch.utils.data.TensorDataset(train_x, train_dataset.targets)
test_ds = torch.utils.data.TensorDataset(test_x, test_dataset.targets)
```

▼ DataLoader

```
batch_size = 32
train_loader = torch.utils.data.DataLoader(
    dataset=train_ds,
    batch_size=batch_size,
    shuffle=True
)

test_loader = torch.utils.data.DataLoader(
    dataset=test_ds,
    batch_size=batch_size,
    shuffle=False # Not necessary!
)
```

▼ Model

```
class Model(nn.Module):
    def __init__(self):
        super(Model, self).__init__()
        self.linear1 = nn.Linear(784, 30) # Flattened MNIST as Input: 28 * 28
        self.relu = nn.ReLU()
        self.selu = nn.SELU()
        self.linear2 = nn.Linear(30, 30)
        self.linear3 = nn.Linear(30, 2) # 2 Classes for The Output: 0-1

    def forward(self, Xa):

        out = self.linear1(Xa)
        out = self.selu(out)
        out = self.linear2(out)
        out = self.selu(out)
        out = self.linear3(out)
        out = self.selu(out)

        return out

# Instantiate the Model
model = Model()
```

▼ Move Model to GPU (Required Condition by Assignment!)

```
device = torch.device("cuda:0" if torch.cuda.is_available() else "cpu")
print(device)
model.to(device)
```

```
cpu
Model(
  (linear1): Linear(in_features=784, out_features=30, bias=True)
  (relu): ReLU()
  (selu): SELU()
  (linear2): Linear(in_features=30, out_features=30, bias=True)
  (linear3): Linear(in_features=30, out_features=2, bias=True)
)
```

▼ Model Summary

```
summary(model, [(1,784)])
```

```
-----
              Layer (type)              Output Shape          Param #
=====
              Linear-1                  [-1, 1, 30]             23,550
              SELU-2                    [-1, 1, 30]              0
              Linear-3                  [-1, 1, 30]             930
              SELU-4                    [-1, 1, 30]              0
              Linear-5                  [-1, 1, 2]               62
              SELU-6                    [-1, 1, 2]               0
=====
Total params: 24,542
Trainable params: 24,542
Non-trainable params: 0
-----
Input size (MB): 0.00
Forward/backward pass size (MB): 0.00
Params size (MB): 0.09
Estimated Total Size (MB): 0.10
-----
```

▼ Loss and Optimizer

```
loss_fn = nn.CrossEntropyLoss()
optimizer = torch.optim.Adam(model.parameters())
```

▼ Training Loop

```
n_epochs = 20
```

```

table_context = tp.TableContext(headers=['epoch', 'Train Acc', 'Train Loss', 'Valid
table_context.__enter__()
for epoch in range(n_epochs):
    train_loss = []
    n_correct = 0.
    n_total = 0.

```

```

for inputs, targets in train_loader:

```

```

    # Move data to GPU
    inputs, targets = inputs.to(device), targets.to(device)

```

```

    # zero the gradient
    optimizer.zero_grad()

```

```

    # forward pass
    output = model(inputs)
    loss = loss_fn(output, targets)

```

```

    # get prediction
    _, prediction = torch.max(output, 1)

```

```

    # update counts
    n_correct += (prediction == targets).sum().item()
    n_total += targets.shape[0]

```

```

    # backward pass and optimize
    loss.backward()
    optimizer.step()

```

```

    train_loss.append(loss.item())

```

```

train_loss = np.mean(train_loss)
train_acc = n_correct / n_total * 100

```

```

test_loss = []
n_correct = 0.
n_total = 0.
tpt, tnt, fpt, fnt = 0, 0, 0, 0
for inputs, targets in test_loader:

```

```

    # Move data to GPU
    inputs, targets = inputs.to(device), targets.to(device)

```

```

    # forward pass
    output = model(inputs)
    loss = loss_fn(output, targets)

```

```

    # get predictions
    _, prediction = torch.max(output, 1)

```

```

    # update counts
    n_correct += (prediction == targets).sum().item()

```

```

n_total += targets.shape[0]

test_loss.append(loss.item())
tp, tn, fp, fn = score(prediction, targets)
tpt += tp
tnt += tn
fpt += fp
fnt += fn

test_loss = np.mean(test_loss)
test_acc = n_correct / n_total * 100
a = (tpt+tnt) / (tpt+tnt +fpt +fnt)
r = tpt / (tpt + fnt)
p = tpt / (tpt + fpt)
f1 = 2 * (p*r)/(p+r)
table_context([epoch+1, train_acc, train_loss, test_acc, test_loss, a, r, p, f1])
table_context.__exit__()

```



epoch	Train Acc	Train Loss	Valid Acc	Valid Loss	V
1	90.043	0.24497	96.83	0.098878	
2	96.978	0.091356	96.99	0.098771	
3	97.54	0.072313	97.83	0.072959	
4	97.878	0.062474	98.02	0.066016	
5	98.103	0.057452	97.29	0.082461	
6	98.317	0.051321	98.09	0.063521	
7	98.402	0.047289	97.48	0.085557	
8	98.535	0.042866	98.17	0.059476	
9	98.55	0.041795	98.32	0.055634	
10	98.732	0.037779	98.1	0.067117	
11	98.762	0.036443	97.77	0.068822	
12	98.843	0.035164	97.98	0.084554	
13	98.942	0.030886	98.19	0.059014	
14	99.002	0.029302	98.18	0.066869	
15	98.982	0.029094	98.15	0.066012	
16	99.092	0.026849	98.13	0.068708	
17	99.077	0.027501	98.16	0.064794	
18	99.165	0.024696	98.31	0.07599	
19	99.13	0.024841	98.26	0.084453	
20	99.225	0.024032	98.4	0.083581	

